

Ekta Shukla - A20567127

Rithika Kavitha Suresh - A20564346

Roger Kewin Samson - A20563057

Jude Rosun - A20564339

Model Selection k-fold cross-validation and bootstrapping

```
In [11]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split, KFold
from sklearn.preprocessing import OneHotEncoder, StandardScaler, PolynomialFeatures
from sklearn.compose import ColumnTransformer
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
from sklearn.utils import resample
```

Load the dataset

```
In [12]: file_path = "flight_prediction.csv" # Replace with the correct file path
data = pd.read_csv(file_path)

# Drop unnecessary columns and clean missing values
data = data.drop(columns=['Unnamed: 0']) # Drop index-like column
data = data.dropna() # Remove rows with missing values

X = data.drop(columns=['price']) # Features
y = data['price'] # Target variable
```

```
In [13]: categorical_features = ['airline', 'source_city', 'departure_time', 'stops',
                                'arrival_time', 'destination_city', 'class']
numerical_features = ['duration', 'days_left']

preprocessor = ColumnTransformer(
```

```

transformers=[
    ('num', StandardScaler(), numerical_features),
    ('cat', OneHotEncoder(handle_unknown='ignore'), categorical_features)
]

# Transform the dataset
X_transformed = preprocessor.fit_transform(X)

```

Split the data into training and testing sets

```

In [14]: X_train, X_test, y_train, y_test = train_test_split(X_transformed, y, test_size=0.2, random_state=42)

X_train_dense = X_train.toarray()
X_test_dense = X_test.toarray()

```

Initialize Linear Regression

```

In [15]: lr = LinearRegression()

# Implementing K-Fold Cross-Validation
kf = KFold(n_splits=10, shuffle=True, random_state=0)
cv_scores = []

for train_index, test_index in kf.split(X_train_dense):
    # Split data manually for each fold
    X_train_kf, X_test_kf = X_train_dense[train_index], X_train_dense[test_index]
    y_train_kf, y_test_kf = y_train.iloc[train_index], y_train.iloc[test_index]

    lr.fit(X_train_kf, y_train_kf)
    y_pred_kf = lr.predict(X_test_kf)

    rss = np.sum((y_test_kf - y_pred_kf)**2)
    tss = np.sum((y_test_kf - np.mean(y_test_kf))**2)
    r2_kf = 1 - (rss / tss)
    cv_scores.append(r2_kf)

# Calculate the mean R^2 score across folds

```

```
cv_mean_r2 = np.mean(cv_scores)
print(f"K-Fold Cross-Validation Mean R2 Score: {cv_mean_r2:.4f}")
```

K-Fold Cross-Validation Mean R² Score: 0.9115

Bootstrapping Implementation

```
In [16]: n_bootstraps = 50 # Number of bootstrap iterations
bootstrap_sample_size = int(0.5 * len(X_train_dense)) # Use 50% of the training data per iteration
bootstrap_scores = []

for _ in range(n_bootstraps):
    indices = np.random.choice(len(X_train_dense), size=bootstrap_sample_size, replace=True)
    X_bootstrap, y_bootstrap = X_train_dense[indices], y_train.iloc[indices]

    # Train and predict
    lr.fit(X_bootstrap, y_bootstrap)
    y_pred_bootstrap = lr.predict(X_test_dense)

    rss = np.sum((y_test - y_pred_bootstrap)**2)
    tss = np.sum((y_test - np.mean(y_test))**2)
    r2_bootstrap = 1 - (rss / tss)
    bootstrap_scores.append(r2_bootstrap)

# Calculate the mean and standard deviation of R2
bootstrap_mean_r2 = np.mean(bootstrap_scores)
bootstrap_std_r2 = np.std(bootstrap_scores)
print(f"Bootstrapping Mean R2 Score: {bootstrap_mean_r2:.4f}")
```

Bootstrapping Mean R² Score: 0.9113

Print final results

```
In [24]: print(f"Final Results:")
print(f" K-Fold Cross-Validation Mean R2: {cv_mean_r2:.4f}")
print(f" Bootstrapping Mean R2: {bootstrap_mean_r2:.4f}")
print(f" Bootstrapping R2 Std Dev: {bootstrap_std_r2:.4f}")
```

Final Results:

K-Fold Cross-Validation Mean R^2 : 0.9115

Bootstrapping Mean R^2 : 0.9113

Bootstrapping R^2 Std Dev: 0.0000

```
In [17]: # AIC
def calculate_aic(n, rss, k):
    return n * np.log(rss / n) + 2 * k
```

```
In [20]: kf_aic_scores = []

for train_index, test_index in kf.split(X_train_dense):
    # Split data manually for each fold
    X_train_kf, X_test_kf = X_train_dense[train_index], X_train_dense[test_index]
    y_train_kf, y_test_kf = y_train.iloc[train_index], y_train.iloc[test_index]
    lr.fit(X_train_kf, y_train_kf)
    y_pred_kf = lr.predict(X_test_kf)

    rss = np.sum((y_test_kf - y_pred_kf)**2)
    tss = np.sum((y_test_kf - np.mean(y_test_kf))**2)
    r2_kf = 1 - (rss / tss)
    cv_scores.append(r2_kf)
    n = len(y_test_kf)
    k = X_train_kf.shape[1] + 1
    aic = calculate_aic(n, rss, k)
    kf_aic_scores.append(aic)
# Print average AIC for K-Fold
mean_aic_kf = np.mean(kf_aic_scores)
print(f"K-Fold Cross-Validation Mean AIC: {mean_aic_kf:.4f}")
```

K-Fold Cross-Validation Mean AIC: 423538.1198

```
In [22]: bootstrap_aic_scores = []

for _ in range(n_bootstraps):
    indices = np.random.choice(len(X_train_dense), size=bootstrap_sample_size, replace=True)
    X_bootstrap, y_bootstrap = X_train_dense[indices], y_train.iloc[indices]
    lr.fit(X_bootstrap, y_bootstrap)
    y_pred_bootstrap = lr.predict(X_test_dense)
    rss = np.sum((y_test - y_pred_bootstrap)**2)
    tss = np.sum((y_test - np.mean(y_test))**2)
    r2_bootstrap = 1 - (rss / tss)
    bootstrap_scores.append(r2_bootstrap)
```

```
n = len(y_test)
k = X_bootstrap.shape[1] + 1
aic = calculate_aic(n, rss, k)
bootstrap_aic_scores.append(aic)

# Print average AIC for Bootstrapping
mean_aic_bootstrap = np.mean(bootstrap_aic_scores)
print(f"Bootstrapping Mean AIC: {mean_aic_bootstrap:.4f}")
```

Bootstrapping Mean AIC: 1058920.7911

In []: